

Real-Time- Stock Market Data Analysis using Kafka

Commands Used:

Login to EC2 Machine

Navigate to the directory where the key is present and then,

- `cd documents/data_engineering-projects-kafka_stock_market_project` (*my key is present in this directory)
- `ssh -i "kafka-stock-market-project.pem" ec2-user@ec2-**-***-***-***.compute-1.amazonaws.com`

-
- `wget https://downloads.apache.org/kafka/3.6.0/kafka_2.12-3.6.0.tgz`
(install kafka)
 - `tar -xvf kafka_2.12-3.6.0.tgz`
(x – extract the zip file,
v- be verbose about it,
f-the next argument is going to be the file name that is to be extracted)

Install Java

- `java -version`
- `wget https://download.oracle.com/java/20/archive/jdk-20_linux-x64_bin.rpm`
- `sudo rpm -ivh jdk-20_linux-x64_bin.rpm`
- `java -version`

-
- `cd kafka_2.12-3.6.0`
-

Start the Zookeeper Server

- `bin/zookeeper-server-start.sh config/zookeeper.properties`

Open another terminal window to start kafka
But first ssh to ec2 machine as done above.

Start Kafka Server

- export KAFKA_HEAP_OPTS="-Xmx256M -Xms128M" ****TO INCREASE THE MEMORY****
- cd kafka_2.12-3.6.0
- bin/kafka-server-start.sh config/server.properties

It points to private server, change server.properties so that it can run in public IP

To do this,

- sudo nano config/server.properties

In this, change ADVERTISED_LISTENERS to public ip of the EC2 instance

Create the Topic

Open another terminal window to create topic and start producer.
But first ssh to ec2 machine as done above.

- cd kafka_2.12-3.6.0
- bin/kafka-topics.sh --create --topic demo_test --bootstrap-server ****_**_**_:9092** --replication-factor 1 --partitions 1

bin/kafka-topics.sh --create --topic demo_testing2 --bootstrap-server {Put the Public IP of your EC2 Instance:9092} --replication-factor 1 --partitions 1

(**Error::** Timed out waiting for a node assignment

Resolution: In EC2->Security->Security Groups-> edit inbound rules -> instead of myIP, select all IP4 -----> This worked for me and the topic got created successfully.)

Start Producer

- `bin/kafka-console-producer.sh --topic demo_test --bootstrap-server **_***_***_***:9092`

Start Consumer

Open another terminal window to create start consumer.
But first ssh to ec2 machine as done above.

- `cd kafka_2.12-3.6.0`
- `bin/kafka-console-consumer.sh --topic demo_test --bootstrap-server **_***_***_***:9092`

So now in

- Terminal 1 – ZooKeeper will be running.**
- Terminal 2 – Kafka will be running.**
- Terminal 3 – Producer will be running.**
- Terminal 4 – Consumer will be running.**